

EpigeneticsBench: Verifiable Evaluation of AI Agents on Epigenomics Analysis

Short-horizon tasks from real epigenomics workflows

Harihara Muralidharan¹, Reema Baskar¹, Soo Hee Lee¹, Kenny Workman¹

¹LatchBio, San Francisco, CA, USA

Correspondence: kenny@latch.bio

ABSTRACT

We introduce EpigeneticsBench, a verifiable benchmark for short-horizon epigenomics analysis. Rather than asking agents to produce open-ended biological interpretations, EpigeneticsBench evaluates whether they can execute focused analysis decisions from realistic workflow states and return deterministically gradable answers. The updated benchmark includes 106 evaluations across CUT&Tag/CUT&RUN, ATAC-seq, ChIP-seq, and DNA methylation workflows. Across 5,088 parsed trajectories from 16 model-harness pairs, no system solves a majority of attempts: GPT-5.5 / Pi leads at 151/318 attempts (47.5%), followed by GPT-5.5 / OpenAI Codex at 135/318 attempts (42.5%) and GPT-5.4 / Pi at 134/318 attempts (42.1%). Difficulty is structured by assay source and workflow convention more than by nominal horizon label, and component-level scoring reveals frequent partial recovery despite endpoint failure. These results suggest that current agents can often locate and manipulate the right biological evidence, but still fail at brittle local decisions involving metric definitions, comparator scope, normalization state, and answer-surface constraints.

Topline Benchmark Performance

We ran the benchmark across 16 model-harness pairs. Pi denotes the Pi terminal coding harness. Passing requires exact recovery of the graded structured answer for an evaluation attempt. The updated manifest contains 5,088 valid result JSONs in a balanced three-replicate design: 106 evaluations for each of 16 model-harness pairs.




Figure 1: Topline benchmark performance. Top: run-level pass rate by model-harness pair across all parsed attempts. Bottom left: pass rate by assay and task family. Bottom right: pass rate against median runtime. Wilson confidence intervals and evaluation-level replicate statistics are reported in Table 2.

Introduction

Epigenomics connects regulatory sequence, chromatin accessibility, histone modification, DNA methylation, and genome organization to cell state and gene regulation. Modern assays can measure these layers at population, single-cell, and spatial resolution, but the resulting analyses are brittle: small choices in genome build, peak filtering, motif background selection, sample aggregation, or denominator definition can change a biological conclusion [1, 2, 3, 4].

AI agents can now inspect files, write analysis code, run command-line tools, and return structured biological summaries. Benchmarks have therefore moved from question answering toward executable scientific tasks. SpatialBench tests localized spatial biology analysis, scBench tests single-cell workflows, and SpatialBench-Long tests recovery of complex scientific conclusions from raw or near-raw spatial data

[5, 6, 7].

Epigenomics needs a complementary short-horizon benchmark. Many consequential failures occur in compact analysis decisions rather than in full-project autonomy: deciding which count layer to use, whether a statistic should be computed per peak, per gene, or per sample, which contrast is biologically appropriate, or how a CpG, motif, or chromatin-state summary should be normalized.

This benchmark evaluates those local empirical decisions. Each task snapshots a realistic analysis state immediately before a target result. The agent must interact with the provided data and return a structured answer that can be graded deterministically. The benchmark is therefore not a test of broad regulatory-biology plausibility. It asks whether an agent can produce the specific verifiable result supported by the data, assay context, and requested answer surface.

Benchmark Design

The benchmark contains 106 evaluations derived from practical epigenomics analysis workflows. Tasks are short-horizon: they isolate one focused analysis decision while preserving enough real data context that a model must inspect files, choose the relevant representation, and compute or verify the requested result. The benchmark follows the SpatialBench convention of withholding the full task set to reduce contamination while reporting aggregate model results and releasing representative examples.

Benchmark Construction Workflow




Figure 2: Benchmark construction workflow. Short-horizon epigenomics evaluations are built by selecting real analysis snapshots, defining a single verifiable target decision, constraining the final answer schema, and using trajectory inspection to identify shortcut-prone or under-specified tasks.

The benchmark is organized around analysis states rather than independent datasets. A single source study or workflow can contribute multiple evaluations that vary the target statistic, comparator, or biological claim. This mirrors how working epigenomics analyses proceed: a scientist repeatedly queries the same assay context to answer adjacent but distinct questions.

Tasks are designed to specify *what* should be recovered rather than *how* it should be recovered. Some procedural evaluations name an operation directly, but most scientific evaluations leave room for the agent to choose reasonable commands, scripts, and summary logic. This prevents the prompt from becoming a hidden solution path while keeping the answer surface verifiable.

The updated inventory contains 97 scientific evaluations and 9 procedural evaluations. Most graders are numeric tolerance checks (97/106 evaluations), with 9 all-of graders for compound structured answers. Component field scores are retained for diagnostics, but run-level pass/fail is the benchmark score. This distinction matters because some grader score fields are not normalized across all evaluations, whereas pass/fail has a consistent interpretation.

Manual trajectory inspection is used in the same spirit as SpatialBench and SpatialBench-Long: not to replace deterministic endpoint grading, but to make failure modes legible. The current metadata marks recurring traps such as wrong comparator or baseline, feature-namespace mismatch, wrong data layer, metric-definition mismatch, and prior-driven biological overinterpretation.

Motivating Example: One Local Decision Can Change

Short-Horizon Epigenomics Chokepoints




Figure 3: Short-horizon epigenomics chokepoints. A representative evaluation snapshots an analysis immediately before a local decision. Agents must identify the relevant assay layer, choose the requested comparison, compute the target statistic over the correct unit of analysis, and return a constrained answer that can be graded deterministically.

The short-horizon setting is easiest to see through a single evaluation. One representative ATAC-seq task begins after peak-level outputs and sample metadata have already been staged. The agent is not asked to reconstruct an entire study; it must choose the correct representation, comparator, denominator, and final answer schema for one constrained empirical result.

This kind of task is small enough to be verifiable but still scientifically fragile. A trajectory can inspect the right files and compute a plausible statistic, yet fail by using the wrong feature namespace, aggregating over the wrong biological unit,

or substituting a nearby biological contrast for the requested one. Conversely, a passing answer requires more than memorized regulatory biology: it requires matching the requested output to the data layer and comparison actually supplied in the task.

This example illustrates why deterministic endpoint grading and diagnostic metadata are both useful. The benchmark score remains binary, because the target answer is a constrained scientific artifact. The failure analysis asks which local decision broke when the endpoint answer was wrong.

Evaluation Inventory

The benchmark covers four assay types and eight task categories. CUT&Tag/CUT&RUN tasks are drawn from zebrafish chromatin workflow snapshots and span chromatin state analysis, QC, peak calling, secondary analysis, genomic annotation, alignment, and visualization. ATAC-seq tasks cover B-ALL and cold-response analyses, with secondary analysis dominating. ChIP-seq tasks include peak calling, chromatin state analysis, QC, genomic annotation, and secondary analysis. Methylation tasks are derived from GSE149608 and GSE149609, covering QC, secondary analysis, genomic annotation, and differential methylation.

Assay type	Evals	Sources	Task categories
CUT&Tag/CUT&RUN	47	Zebrafish chromatin workflow snapshots	Chromatin state (12), QC (10), peak calling (9), secondary analysis (7), genomic annotation (5), alignment (3), visualization (1)
ATAC-seq	24	B-ALL ATAC, cold-response ATAC	Secondary analysis (17), chromatin state (2), peak calling (2), QC (2), genomic annotation (1)
ChIP-seq	10	ChIP-seq workflow snapshots	Peak calling (3), chromatin state (2), QC (2), secondary analysis (2), genomic annotation (1)
Methylation-seq	25	GSE149608, GSE149609	QC (8), secondary analysis (8), genomic annotation (5), differential methylation (4)

Table 1: Evaluation inventory. The benchmark groups evaluations by assay type while varying the source workflow, task category, and expected answer surface within each group.

Across all 106 evaluations, 34 are secondary analysis tasks, 22 are QC, 16 are chromatin state analysis, 14 are peak calling, 12 are genomic annotation, 4 are differential methylation, 3 are alignment, and 1 is visualization. The updated result manifest contains 5,088 valid trajectories from 16 model-harness pairs, with exactly 318 attempts per model-harness pair.

Results

Endpoint grading identifies a leading tier rather than a fully ordered leaderboard

Across 16 model-harness pairs and 5,088 valid trajectories, deterministic final-answer grading produced 1,672 passes (32.9%). Success was concentrated but not cleanly separated across the leading systems. GPT-5.5 / Pi achieved the highest pass rate, passing 151/318 attempts (47.5%; Wil-

son 95% confidence interval (CI), 42.1–53.0). GPT-5.5 / OpenAI Codex followed at 135/318 attempts (42.5%; 37.1–47.9), while GPT-5.4 / Pi reached 134/318 attempts (42.1%; 36.8–47.6). Claude Opus 4.8 Max / Pi and Claude Opus 4.8 / Claude Code both passed 126/318 attempts (39.6%).

The updated manifest is balanced: each model-harness pair contributes three rounds for each of 106 evaluations. The practical interpretation is that GPT-5.5 / Pi is the descriptive leader, while the nearby systems form a leading tier whose Wilson intervals overlap substantially.




Figure 4: Endpoint success across model-harness pairs. Left: run-level pass rate by model-harness pair across all 5,088 parsed attempts with Wilson confidence intervals. Right: evaluation-level replicate statistics, reporting the fraction of attempted evaluations with at least one passing attempt, a majority of passing attempts, or all available attempts passing.

Model / harness	Pass runs	Pass %	Wilson 95% CI	Any	Majority	All	Cost	Median min
GPT-5.5 / Pi	151/318	47.5	42.1–53.0	61/106	52/106	38/106	\$1.284	13.4
GPT-5.5 / OpenAI Codex	135/318	42.5	37.1–47.9	57/106	45/106	33/106	–	7.2
GPT-5.4 / Pi	134/318	42.1	36.8–47.6	58/106	49/106	27/106	\$0.924	15.3
Claude Opus 4.8 Max / Pi	126/318	39.6	34.4–45.1	50/106	41/106	35/106	–	11.8
Claude Opus 4.8 / Claude Code	126/318	39.6	34.4–45.1	51/106	44/106	31/106	–	10.2
GPT-5.4 / OpenAI Codex	122/318	38.4	33.2–43.8	54/106	38/106	30/106	–	8.0
Claude Opus 4.7 / Claude Code	109/318	34.3	29.3–39.7	48/106	35/106	26/106	–	6.1
Claude Opus 4.7 / Pi	105/318	33.0	28.1–38.4	45/106	36/106	24/106	–	6.9
Claude Opus 4.6 / Pi	104/318	32.7	27.8–38.0	46/106	34/106	24/106	\$0.620	7.2
Gemini 3.5 Flash / Pi	98/318	30.8	26.0–36.1	46/106	33/106	19/106	\$1.587	14.8
Claude Opus 4.6 / Claude Code	97/318	30.5	25.7–35.8	50/106	30/106	17/106	–	8.0
Gemini 3.1 Pro / Pi	86/318	27.0	22.5–32.2	41/106	29/106	16/106	\$0.768	12.3
Claude Sonnet 4.6 / Pi	81/318	25.5	21.0–30.5	38/106	26/106	17/106	\$0.609	7.4
Kimi K2P6 / Pi	79/318	24.8	20.4–29.9	44/106	21/106	14/106	\$0.460	12.0
Grok 4.20 reasoning / Pi	76/318	23.9	19.5–28.9	40/106	25/106	11/106	\$0.280	4.7
Grok 4.3 / Pi	43/318	13.5	10.2–17.7	21/106	13/106	9/106	\$0.033	0.7

Table 2: Deterministic final-answer performance. Run-level pass rates are computed over the balanced updated manifest. Any, Majority, and All report evaluation-level replicate statistics over 106 evaluations for each model-harness pair. Cost is shown only when it is present for all attempts from that model-harness pair; dashes indicate incomplete or unavailable cost metadata. Median minutes summarizes the recorded result duration.

Assay source stratifies difficulty more clearly than horizon labels

Performance varied more by assay source and workflow provenance than by the coarse horizon label. ChIP-seq tasks passed at 241/480 attempts (50.2%), followed by CUT&Tag/CUT&RUN at 768/2,256 (34.0%), methylation at 400/1,200 (33.3%), and ATAC-seq at 263/1,152 (22.8%). Source-level summaries exposed a sharper split within these assay labels: GSE149608 methylation tasks passed at 388/1,152 attempts (33.7%), whereas the single GSE149609 task passed at 12/48 (25.0%). Cold-response ATAC tasks passed at 147/576 (25.5%), while B-ALL ATAC-prefix tasks passed at 116/576 (20.1%).

Task family alone was less informative. Secondary analysis

tasks passed at approximately one-third of attempts, while QC tasks passed at a similar rate. However, the same nominal task family behaved differently across assay contexts: CUT&Tag/CUT&RUN secondary analysis tasks passed at higher rates than ATAC-seq secondary analysis tasks. Agents are therefore not simply failing broad task types; they are sensitive to the representation, comparator, and denominator conventions of particular epigenomics workflows.

The nominal horizon label (short vs. long) is almost entirely confounded with assay type: all CUT&Tag/CUT&RUN and ATAC-seq evaluations are short horizon, while all methylation evaluations are long horizon. Any apparent difference in pass rate between horizon groups is therefore a source-composition effect rather than evidence that task duration drives difficulty (Figure 5).




Figure 5: Performance varies by assay type and task category. Left: pass rate by assay type with run-level Wilson intervals. Middle: assay-by-task-category pass-rate heatmap; cell text reports pass rate. Right: evaluation count by assay type, showing the unbalanced inventory composition.

Replicate statistics distinguish occasional recovery from stable recovery

Run-level pass rates are useful for ranking systems, but evaluation-level replicate statistics better capture reliability. GPT-5.5 / Pi passed at least one attempt on 61/106 evaluations, passed a majority of attempts on 52/106, and passed all three attempts on 38/106. GPT-5.5 / OpenAI Codex showed a similar drop from 57/106 any-pass evaluations to 45/106 majority-pass and 33/106 all-pass evaluations. In contrast, Grok 4.3 / Pi passed at least one attempt on 21/106 evaluations and a majority on 13/106.

This gap matters because short-horizon tasks are not purely deterministic from the agent’s perspective. A model can know the relevant assay logic but choose the wrong field name, denominator, or comparator in one replicate. Conversely, a single passing replicate does not imply stable operational competence. We therefore treat any-pass rates as evidence of reachable capability and majority/all-pass rates as stronger evidence of repeatable recovery.

Hard evaluations often reflect component conjunction failures

Evaluation difficulty was highly skewed. Across 106 evaluations, 21 had no passing attempts and 18 more had pass rates below 10%, whereas no evaluation was passed by every available attempt and only three exceeded 90% pass rate. The hardest evaluations were not uniformly cases of no analysis progress: 17 of 21 zero-pass evaluations had nonzero mean grader score, and 9 of 21 had field-level pass rates of at least 50%.

Field scores show that many failures occur at the final conjunction of components. Across all 5,088 runs, endpoint pass/fail was exactly equivalent to all component fields passing. The weighted field-level pass rate was 69.6%, while endpoint pass rate was 32.9%. This gap was largest for ATAC-seq secondary analysis, CUT&Tag/CUT&RUN chromatin state analysis, ATAC-seq QC, CUT&Tag/CUT&RUN peak calling, and methylation-seq genomic annotation tasks. These results indicate that many agents recover substantial local structure but miss one required field, comparator, denominator, or answer-surface constraint.




Figure 6: Endpoint failures often contain substantial partial progress. Top: endpoint pass rate and field-level pass rate by assay and task family (dumbbell chart). Bottom left: evaluation-level difficulty spectrum, with endpoint and field-level pass rate for each evaluation. Bottom right: annotated failure-mode pass rates. Field scores and failure annotations are diagnostics, not replacement benchmark scores.

Failure annotations localize errors to concrete analysis decisions

Failure annotations support the component-conjunction interpretation. The most frequent annotated families were metric/statistic mismatch, normalization or representation mismatch, biological-interpretation errors, subset or comparator errors, and threshold-resolution instability. These categories match common epigenomics failure modes: reaching the right data object but summarizing the wrong layer, computing a plausible statistic with the wrong denominator, or substituting a nearby biological contrast for the requested comparison.

Annotation coverage is incomplete. Only 25/106 evaluations currently have failure-mode labels, including 10/21 zero-pass evaluations. Failure modes are also evaluation-level annotations rather than per-run observed labels, and co-annotated evaluations appear in multiple rows of the failure-mode summary. We therefore treat the taxonomy as a diagnostic sample of recurring traps rather than a complete decomposition of benchmark error.

Runtime failures are not the dominant error mode

Among parsed result JSONs, no trajectory reports a timeout or detected out-of-memory event. Recorded runtimes nevertheless vary widely: the median duration is 506 seconds and the 90th percentile is 2,351 seconds. Grok 4.3 / Pi is very fast, with a median duration of 41 seconds and a 13.5% pass rate. GPT-5.5 / Pi is slower, with a median duration of 802 seconds and the highest pass rate. These measurements should be interpreted as operational diagnostics rather than biological capability scores. Cost metadata are available for 2,875/5,088 attempts and are complete for only a subset of model-harness pairs, so they should not be used for a global cost-efficiency ranking in the current manuscript.

Discussion

The short-horizon epigenomics benchmark occupies a middle ground between localized workflow execution and long-horizon scientific reasoning. It is narrower than SpatialBench-Long because each task targets a focused result, but it is more empirical than biological question answering because the answer depends on interacting with real assay outputs.

The main result is not that agents fail universally. Leading systems recover roughly half of available endpoint answers, and many failures show substantial component-level progress. The limitation is reliability: even strong model-harness pairs remain sensitive to representation, metric definition, and comparator choice. These are precisely the local decisions that determine whether an epigenomics analysis is trustworthy.

The benchmark has several limitations. The updated manifest is balanced across model-harness pairs, but the task inventory is not balanced across assay families: CUT&Tag/CUT&RUN (47 evaluations) and methylation-seq (25) contribute more evaluations than ChIP-seq (10), and secondary analysis plus QC tasks dominate the inventory. The nominal horizon label is fully confounded with assay type and should not be interpreted independently. Finally, deterministic graders intentionally constrain the answer surface; they are useful for measurement, but they do not capture every scientifically valid analysis path.

Methods

Evaluation construction

Candidate evaluations were derived from real epigenomics analysis workflows. An evaluation was retained when the target result could be reproduced from the provided data, expressed through a constrained final answer, and graded deterministically. Tasks were revised or removed when prompts over-specified the method, when reasonable analysis choices produced incompatible answers, or when the grader could not distinguish a supported result from plausible but unsupported biological output.

Agent runs and aggregation

Each model-harness pair was run with up to three replicates per evaluation. The updated manifest used here is `epibench_manifest_updated.csv`; the full SHA-256 hash is recorded in the derived summary facts table. All 5,088 listed result JSONs were downloaded and parsed successfully. The manifest contains a complete three-replicate design for 106 evaluations and 16 model-harness pairs. Failed, malformed, or nonpassing endpoint grader outputs are counted as nonpassing attempts when present; no downloaded result JSON was missing or invalid in this analysis.

Statistical analysis

The primary metric is run-level pass rate. Wilson binomial confidence intervals summarize run-level uncertainty and do not account for evaluation-level clustering. We therefore also report evaluation-level replicate summaries: whether a model-harness pair passed any attempt, a majority of attempts, or all three attempts for an evaluation. Because the updated manifest is balanced, run-level and balanced-design summaries are identical. Component field scores and failure-mode labels are reported as diagnostics rather than benchmark scores.

Data availability

Results files, public example evaluations, and representative trajectories will be made available at <https://github.com/latchbio/epigenetics-short-horizon-benchmark>. The derived summary tables used for this manuscript are in `analysis_outputs/epi_manifest_updated_summary/`.

References

- [1] ENCODE Project Consortium. An integrated encyclopedia of DNA elements in the human genome. *Nature* 489, 57–74 (2012).
- [2] Roadmap Epigenomics Consortium. Integrative analysis of 111 reference human epigenomes. *Nature* 518, 317–330 (2015).
- [3] Buenrostro, J. D. et al. Transposition of native chromatin for fast and sensitive epigenomic profiling of open chromatin, DNA-binding proteins and nucleosome position. *Nature Methods* 10, 1213–1218 (2013).
- [4] Kaya-Okur, H. S. et al. CUT&Tag for efficient epigenomic profiling of small samples and single cells. *Nature Communications* 10, 1930 (2019).
- [5] Workman, K., Yang, Z., Muralidharan, H. & Le, H. Spatial-Bench: Can agents analyze real-world spatial biology data? *arXiv* 2512.21907 (2025).
- [6] Workman, K. et al. scBench: Verifiable benchmarking of single-cell analysis agents. Manuscript in preparation (2026).
- [7] Diks, I., Muralidharan, H., Proctor, T. & Workman, K. Verifiable benchmarking of long-horizon spatial biology. *arXiv* 2605.28065 (2026).